

# diabetes-prediction-system

December 28, 2023

```
[1]: #Import
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%pip install xgboost
```

Requirement already satisfied: xgboost in d:\users\mathi\anaconda3\lib\site-packages (2.0.3)

Requirement already satisfied: numpy in d:\users\mathi\anaconda3\lib\site-packages (from xgboost) (1.24.3)

Requirement already satisfied: scipy in d:\users\mathi\anaconda3\lib\site-packages (from xgboost) (1.11.1)

Note: you may need to restart the kernel to use updated packages.

```
[8]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import \
    accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
```

```
[3]: data=pd.read_csv("D:\power bi\Meriskill\diabetes.csv")
data
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	


766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[5]: sns.heatmap(data.isnull())
```

```
[5]: <Axes: >
```



```
[11]: correlation = data.corr()
      print(correlation)
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[10]: sns.heatmap(data.corr())
```

```
[10]: <Axes: >
```



```
[14]: #preparing the data
X=data.drop("Outcome",axis=1)
Y=data['Outcome']
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.
↪2,random_state=42)
```

```
[16]: #standardize the feature
scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)
```

```
[20]: #train a random classifier
model=RandomForestClassifier(random_state=42)
model.fit(X_train_scaled,Y_train)
```

```
[20]: RandomForestClassifier(random_state=42)
```

```
[17]: #evaluation
Accuracy=accuracy_score(Prediction,Y_test)
print(Accuracy)
```

0.7272727272727273

```
[8]: #display the first few rows
print(data.head())
#checking for missing values
print(data.isnull().sum())
#statistical summary
print(data.describe())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0


dtype: int64

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000

mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[13]: #distribution of target variable
sns.countplot(x='Outcome',data=data)
plt.title('Distribution of outcome (0:no diabetes,1:diabetes)')
plt.show()
```



```
[23]: #additional evaluation metrics
print('\n classification report:')
print(classification_report(Y_test,predictions))
```

```
classification report:
      precision    recall  f1-score   support
```

0	0.79	0.78	0.78	99
1	0.61	0.62	0.61	55
accuracy			0.72	154
macro avg	0.70	0.70	0.70	154
weighted avg	0.72	0.72	0.72	154

```
[24]: print('\nConfusion matrix:')
print(confusion_matrix(Y_test,predictions))
```

```
Confusion matrix:
[[77 22]
 [21 34]]
```

```
[29]: sns.set(style="whitegrid")
#1.Pregnancies
plt.figure(figsize=(10,6))
sns.countplot(x='Pregnancies',data=data)
plt.title('No.of Pregnancies')
plt.show()
#2.Glucose
plt.figure(figsize=(10,6))
sns.histplot(x='Glucose',data=data,kde=True)
plt.title('Distribution of Glucose Levels')
plt.show()
#3.Blood Pressure
plt.figure(figsize=(10,6))
sns.histplot(x='BloodPressure',data=data,kde=True)
plt.title('Distribution of Blood Pressure')
plt.show()
#4.Skin THickness
plt.figure(figsize=(10,6))
sns.histplot(x='SkinThickness',data=data,kde=True)
plt.title('Distribution of Skin Thickness')
plt.show()
#5.Insulin
plt.figure(figsize=(10,6))
sns.histplot(x='Insulin',data=data,kde=True)
plt.title('Distribution of Insulin Levels')
plt.show()
#6.BMI
plt.figure(figsize=(10,6))
sns.histplot(x='BMI',data=data,kde=True)
plt.title('Distribution of BMI')
plt.show()
```




```


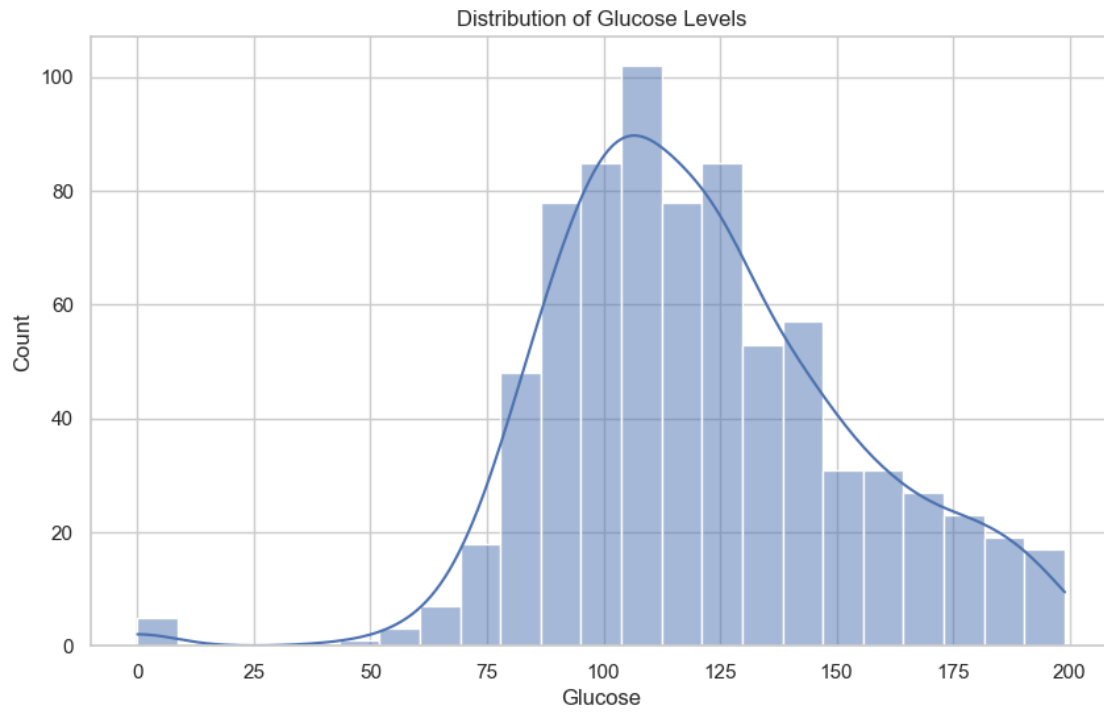
#7.Diabetes Pedigree Function
plt.figure(figsize=(10,6))
sns.histplot(x='DiabetesPedigreeFunction',data=data,kde=True)
plt.title('Distribution of Diabetes Pedigree Function Scores')
plt.show()

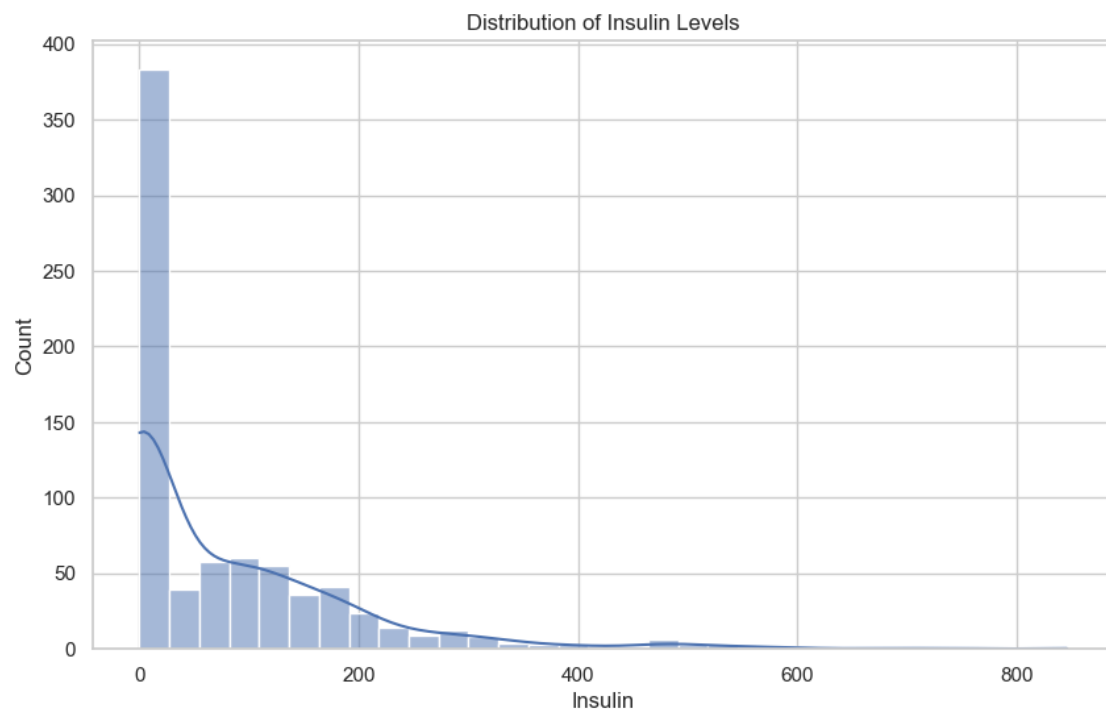

#8.Age
plt.figure(figsize=(10,6))
sns.histplot(x='Age',data=data,kde=True)
plt.title('Distribution of Age')
plt.show()

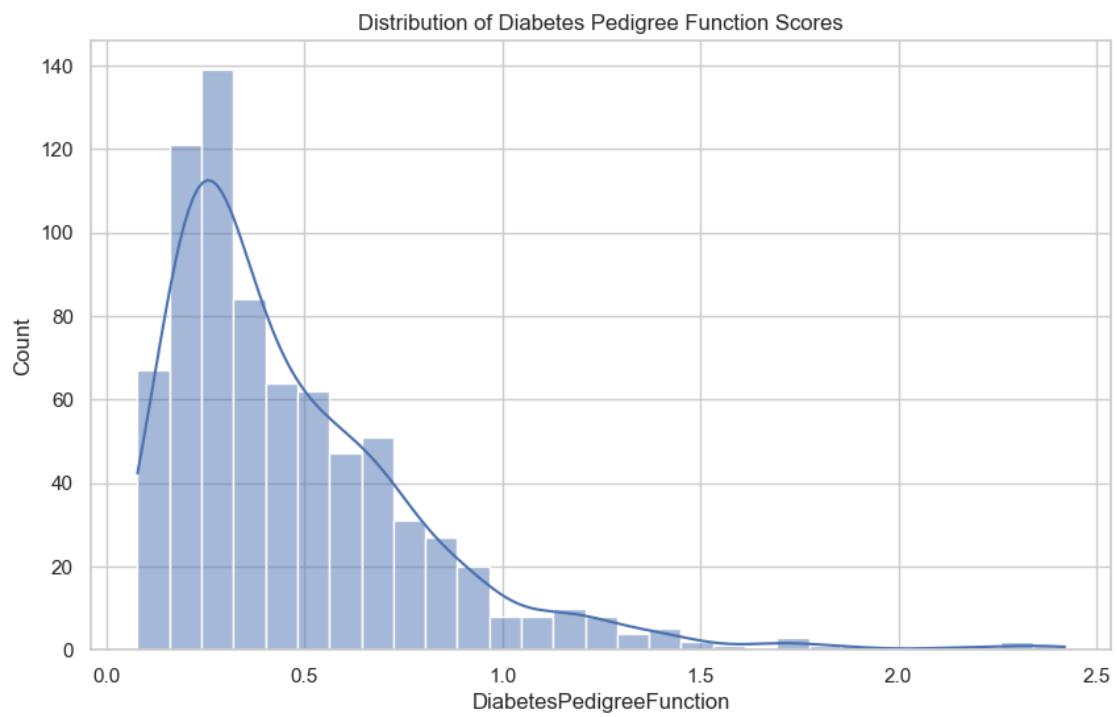
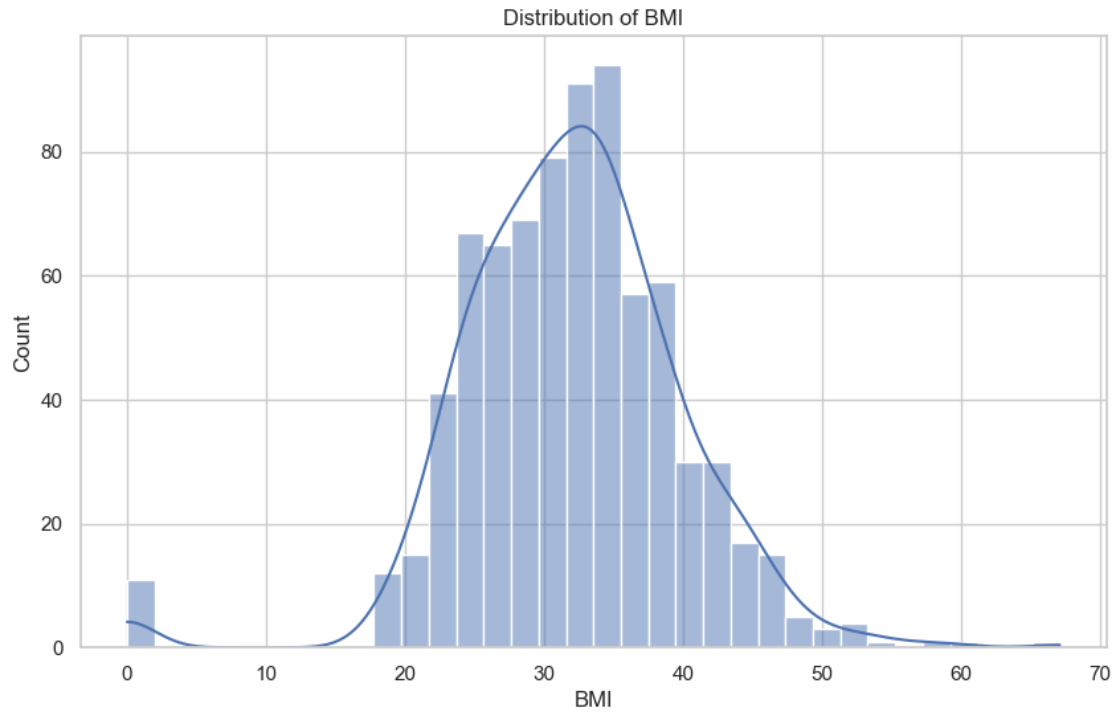
#9.Outcome
plt.figure(figsize=(10,6))
sns.countplot(x='Outcome',data=data)
plt.title('Distribution of Outcome 1: Diabetes (0: No Diabetes)')
plt.show()



```



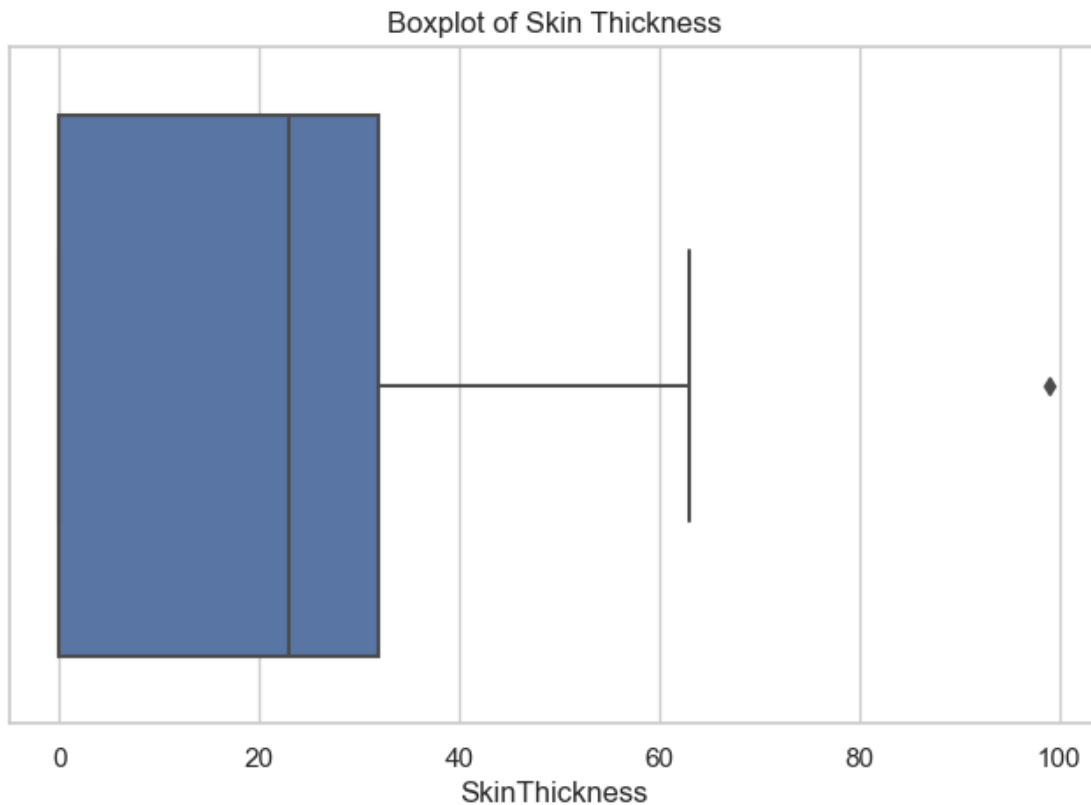








```
[31]: #Boxplot of skin thickness
plt.figure(figsize=(8,5))
sns.boxplot(x='SkinThickness',data=data)
plt.title('Boxplot of Skin Thickness')
plt.show()
```

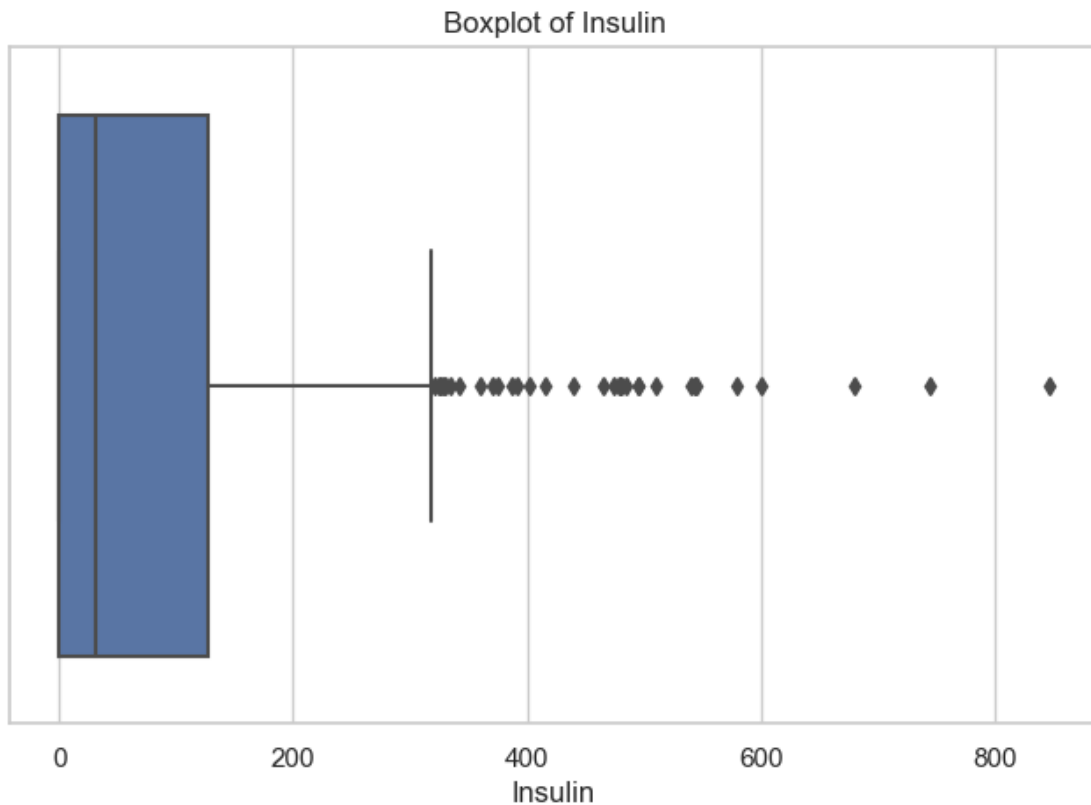


```
[32]: #outliers in skin thickness
Q1_skin_thickness = data['SkinThickness'].quantile(0.25)
Q3_skin_thickness = data['SkinThickness'].quantile(0.75)
IQR_skin_thickness = Q3_skin_thickness - Q1_skin_thickness
lower_bound_skin_thickness = Q1_skin_thickness - 1.5 * IQR_skin_thickness
upper_bound_skin_thickness = Q3_skin_thickness + 1.5 * IQR_skin_thickness
outliers_skin_thickness = data[(data['SkinThickness'] <
    ↳ lower_bound_skin_thickness) | (data['SkinThickness'] >
    ↳ upper_bound_skin_thickness)]

print(f'Number of outliers in SkinThickness: {len(outliers_skin_thickness)}')
```

Number of outliers in SkinThickness: 1

```
[33]: #Boxplot of insulin
plt.figure(figsize=(8,5))
sns.boxplot(x='Insulin',data=data)
plt.title('Boxplot of Insulin')
plt.show()
```



```
[35]: #outliners in insulin
Q1_insulin = data['Insulin'].quantile(0.25)
Q3_insulin = data['Insulin'].quantile(0.75)
IQR_insulin = Q3_insulin - Q1_insulin
lower_bound_insulin = Q1_insulin - 1.5 * IQR_insulin
upper_bound_insulin = Q3_insulin + 1.5 * IQR_insulin
outliers_insulin = data[(data['Insulin'] < lower_bound_insulin) |
    ↪(data['Insulin'] > upper_bound_insulin)]

print(f'Number of outliers in Insulin: {len(outliers_insulin)}')
```

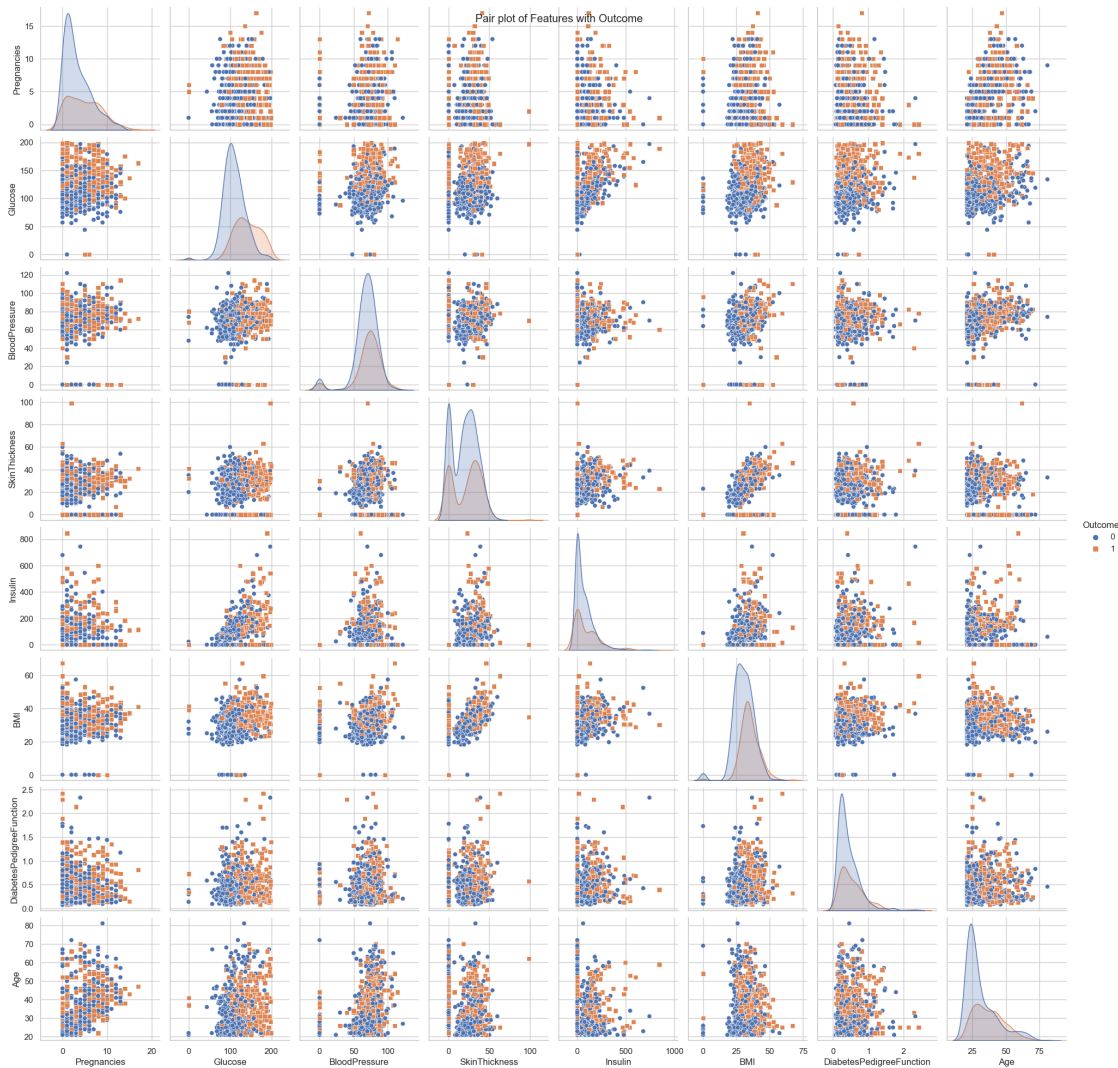
Number of outliers in Insulin: 34

```
[36]: #visualization between features and pair plot
sns.pairplot(data,hue='Outcome',diag_kind='kde',markers=["o","s"])
```

```
plt.suptitle('Pair plot of Features with Outcome')
plt.show()
```


D:\Users\mathi\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:  
The figure layout has changed to tight

```
self._figure.tight_layout(*args, **kwargs)
```

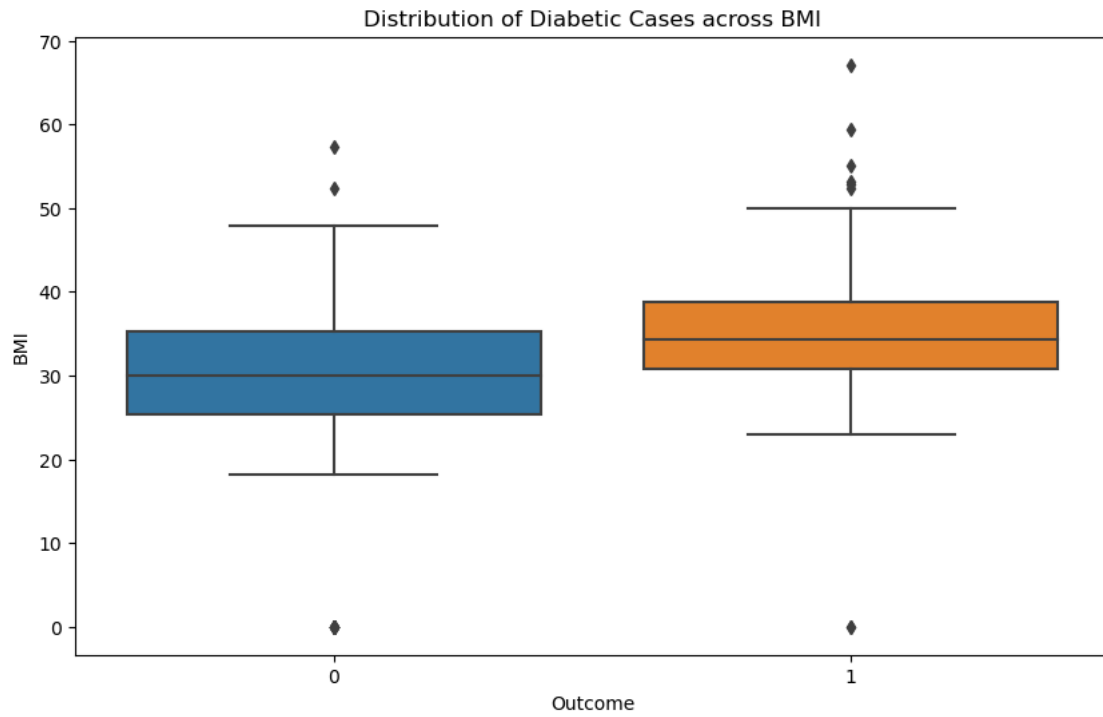


```
[12]: #Distribution of Age
plt.figure(figsize=(10, 6))plt.figure(figsize=(10, 6))
sns.histplot(x='Age', data=data, kde=True)
plt.title('Distribution of Age')
plt.show()
```

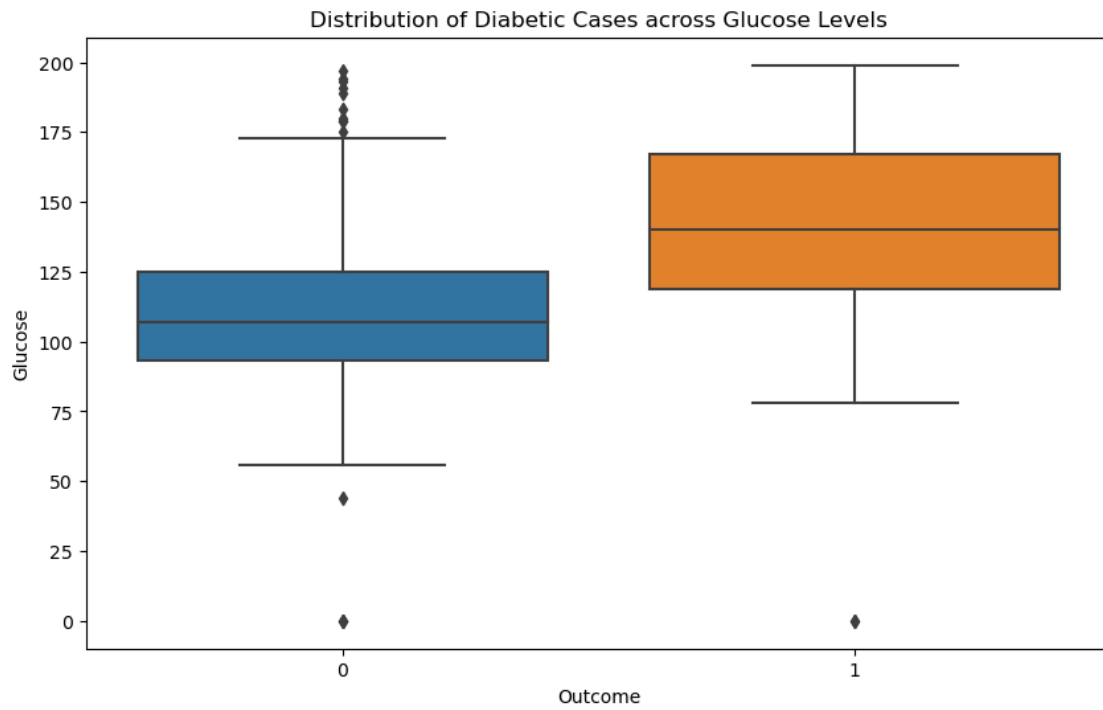




```
[13]: #Distribution of cases accross BMI
plt.figure(figsize=(10, 6))plt.figure(figsize=(10, 6))
sns.boxplot(x='Outcome', y='BMI', data=data)
plt.title('Distribution of Diabetic Cases across BMI')
plt.show()
```



```
[14]: #Distribution of cases accross glucose level  
plt.figure(figsize=(10, 6))  
sns.boxplot(x='Outcome', y='Glucose', data=data)  
plt.title('Distribution of Diabetic Cases across Glucose Levels')  
plt.show()
```



[ ]: